



# Introduction to Applied Machine Learning

BUMIC + DSC React Series


**BOSTON UNIVERSITY**  
**MACHINE INTELLIGENCE**  
**COMMUNITY**

Darcy  
03/09/2021

# Applications of Deep Learning

## 1. Cool things using deep learning

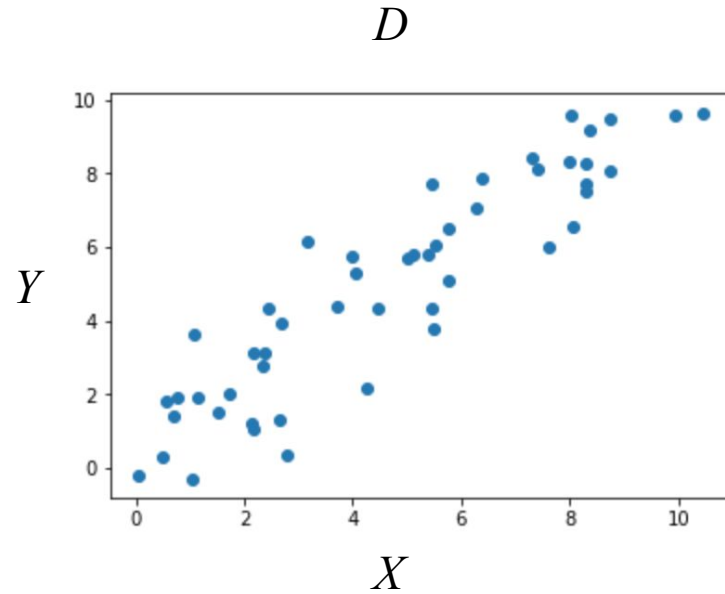
- a. Computer Vision
  - i. Tesla recognizing items on a street
- b. Text generation
  - i. OpenAI GPT3 can solve almost any language task in a few examples
- c. Reinforcement Learning
  - i. Can play Atari games, Board games, Real Time Strategy games
  - ii. Robotic control
- d. Many more...



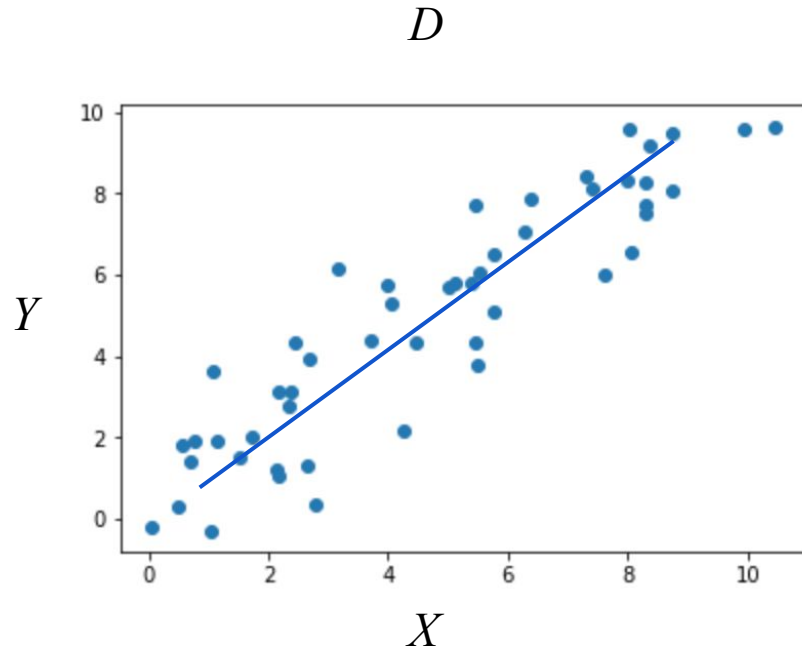


# Learning from data

We have some data  $D$



Make an assumption about  $D$

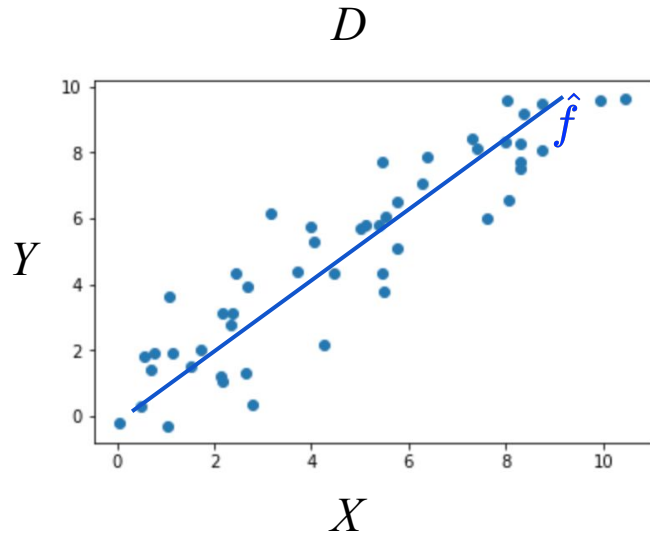


$$y = b + mx$$

$$\hat{f} = \theta_0 + \theta_1 x$$

# What is learning?

The approximation of some unknown function  $f$  based on some data  $D$ .



$$f : X \rightarrow Y$$

$$\hat{f} = \theta_0 + \theta_1 x$$


*How do we set the parameters?*

*How do we know what assumptions to make?*



# Intro to Deep Learning

# What is Deep Learning




Deep learning is a subset of machine learning




# What is Deep Learning

*Deep learning learns from data using a class of functions known as Neural Networks*




*A neural network maps an input to an output*

# Biological Neuron vs. Artificial Neuron




A cartoon drawing of a biological neuron (left) and its mathematical model (right).




# What is a Neural Network?

**Simple Neural Network**



**Deep Learning Neural Network**



● Input Layer

● Hidden Layer

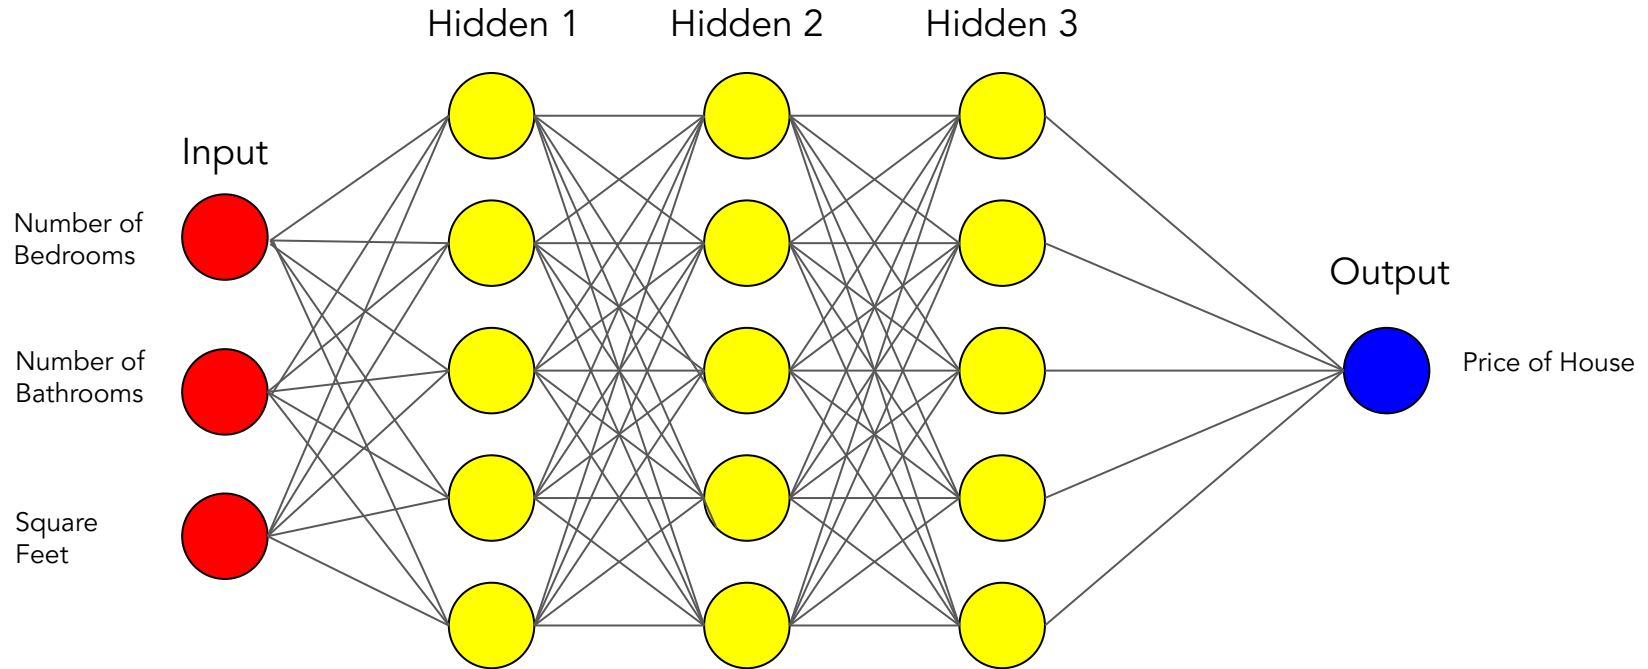
● Output Layer



# Steps to Train a NN

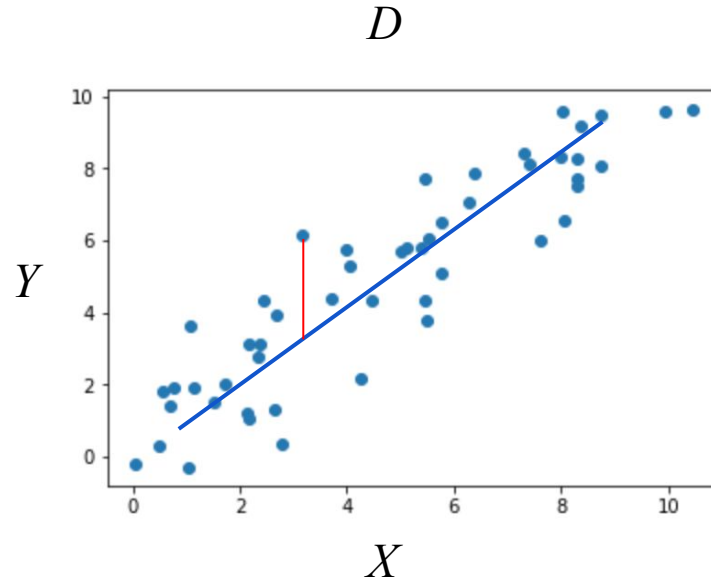

# Forward propagation

Push example through the network to get a predicted output




# Compute the cost

Calculate difference between predicted output and actual data



# Compute the cost

Calculate difference between predicted output and actual data





$$J(\theta) = \frac{1}{2m} \sum_i^m (y_i - \hat{y}_i)^2$$

Where  $i$  is the  $i$ th training example and  $m$  is the number of training examples



# Backward propagation - "Update"

Push back the derivative of the error and apply to each weight, such that next time it will result in a lower error



<https://hmkcode.github.io/ai/backpropagation-step-by-step/>








# Convolutional Neural Networks

# Image Data

- Images are commonly represented in code as a 3D array of pixels. Here, we notice 3 represents RGB values
- In vanilla neural networks, we would simply flatten this 3D array into a 3072 length vector. However, by doing this, we lose spatial correlation between pixels close to other pixels


32x32x3 image



# Image Data

- In 2012 a paper called AlexNet out competed state of the art image classification models through the usage of kernels (also called filters)


32x32x3 image




# Kernel

- Kernel: a small matrix used for feature detection on an image
  - Also called a filter
- Usage
  - Superimpose the kernel over a section of an image
  - Do element-wise multiplication between the weights in the kernel and the values in the image
  - Record the sum of the multiplications

32x32x3 image



5x5x3



# Example Convolution

Example: Multiply the 5x5 image by a 3x3 kernel with weights:

1 0 1  
0 1 0  
1 0 1

The output?  
Sum of weight times  
part of image to a  
single number.

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved  
Feature

# Kernel example

6	3	2
4	3	1
3	5	5

Section of  
an image

\*

0	1	0
1	2	1
0	1	0

Kernel

= sum

6*0	3*1	2*0
4*1	3*2	1*1
3*0	5*1	5*0

=

19

## Kernel example (cont.)

3	3	1
4	6	5
3	5	2

Section of  
an image

\*

0	1	0
1	2	1
0	1	0

Kernel

= sum

$3*0$	$3*1$	$1*0$
$4*1$	$6*2$	$5*1$
$3*0$	$5*1$	$2*0$

=

23
----

This image section contains the same values as before, but they have been rearranged, resulting in a greater activation with this kernel

# Example Convolution

- Note that the output is smaller than the input
- This can be prevented by using padding around the edges of the image.

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image


4		

Convolved  
Feature



# Padding

- Before padding:
  - 7x7 input, 3x3 filter creating a 5x5 sized output
- After padding:
  - 9x9 input, 3x3 filter creating a 7x7 sized output which maintains the same size as our input
- Edges and corners aren't as accurate but in practice this works well enough



# Stride

- Here, the kernel is moving one pixel at a time ("stride" = 1)
- The kernel can move by more than one pixel at a time
- $\text{Size} = (N - F) / \text{Stride} + 1$

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		


Convolved  
Feature

# Stride


- Increasing stride decreases the size of the output
- Here, stride = 2
- $(N - F) / \text{Stride} + 1$

$$(7 - 3) / 2 + 1 = 3$$

7 x 7 Input Volume




3 x 3 Output Volume




# Dimensionality Practice

- What would be the output size of a 5x5x3 filter with a 32x32x3 image and a stride of 1?
- $(N - F) / \text{Stride} + 1$

32x32x3 image




5x5x3



# Dimensionality Practice

- $(32 - 5) / 1 + 1 = 28$
- Now let's say we had a stride of 2,
  - $(32 - 5) / 2 + 1 = 14.5$
  - Fractional size means the filter hangs off the input
  - We wouldn't use this stride value consequently




# Intentionally Shrinking Output Size

- Now, let's say you want to shrink your outputs (which are inputs to the next layer) to reduce operations.
- You can do this by either increasing the stride
  - $(N - F) / \text{Stride} + 1$
- Alternatively, you can use a pooling layer



# Conv Layer Output

- Use multiple kernels for multiple activation maps
- In this example, we have 6 activation maps each created through a different filter with its own set of weights and biases






Additional Layers



# Pooling Layers



Convolved  
feature



Pooled  
feature

- Limitation of output of Convolutional Layers:
  - Record the precise position of features in the input
  - Small movements in the position of the feature in the input image will result in a different feature map
- Solution: Pooling Layers
  - Lower resolution version of input is created with large and important structure elements preserved
  - Reduces the computational cost by reducing the number of parameters to learn


# Max Pooling

Extracts the sharpest features of an image, making it more general



Input (4 x 4)

255	0	0	0
0	255	0	0
0	0	255	0
0	0	0	255




Output (2 x 2)

255	0
0	255


# Average Pooling

Takes average feature of an image, minimize overfitting



Input (4 x 4)

0	255	255	255
255	0	255	255
255	255	0	255
255	255	255	0




Output (2 x 2)

128	255
255	128


# Dropout

## 1. First, what is overfitting?


- Overfitting is when the neural network corresponds too closely to the dataset, and cannot be generalized. This tends to happen when a model is excessively complex relative to the data
- Conversely, underfitting is when the network cannot capture the underlying trend of the dataset which may happen if your network is not complicated enough.



Underfitted



Good Fit/Robust




Overfitted


# Dropout - How can we solve overfitting?

## 1. Training phase

- Each weight has a probability  $p$  that they will be multiplied by zero (dropped). This probability is often set to 0.5, which is considered to be close to optimal for a wide range of networks and tasks
- This has the effect of removing random connections between activations effectively creating a new network/outlook on the data per each train set



(a) Standard Neural Net



(b) After applying dropout.

# Dropout - How can we solve overfitting?

## 2. Post Train

- a. After training weights will be abnormally high as they were adjusted assuming only  $(1-p)$  percent of the weights would be summed together and used.
- b. To fix this we normalize weights to lower the expectation of each weight. We do this by scaling each weight by  $1/p$
- c. "This makes sure that for each unit, the expected output from it under random dropout will be the same as the output during pretraining." ~Dropout: A Simple Way to Prevent Neural Networks from Overfitting
  - i. <http://www.jmlr.org/papers/volume15/srivastava14a/srivastava14a.pdf>



# Convolutional Neural Network

